

Stat 292 - Recitation 12

Probability Distributions & Exceptions and Timings

Orçun Oltulu

16 / 6 / 2021

Probability Distributions in R

Every distribution that R handles has four functions. There is a root name, for example, the root name for the normal distribution is `norm`. This root is prefixed by one of the letters;

- **p** for “probability”, the cumulative distribution function (cdf)
- **q** for “quantile”, the inverse cdf.
- **d** for “density”, the density function (pmf or pdf)
- **r** for “random”, a random variable having the specified distribution

For the normal distribution, these functions are `pnorm`, `qnorm`, `dnorm`, and `rnorm`. For the binomial distribution, these functions are `pbinom`, `qbinom`, `dbinom`, and `rbinom`. And so forth.

Exercise 1:

Bob makes 60% of his free-throw attempts. If he shoots 12 free throws,

Part A:

What is the probability that he makes exactly 10? $P(X = 10) = ?$

```
dbinom(10, size = 12, prob = .6)
```

```
## [1] 0.06385228
```

Part B:

What is the probability that he makes more than 9? $P(X > 9) = ?$

```
1 - pbinom(9, size = 12, prob = .6)
```

```
## [1] 0.08344332
```

```
# alternatively,  
# pbinom(9, size = 12, prob = .6, lower.tail = F)
```

Part C:

Find the 10th quantile of Bob's free-throw attempts.

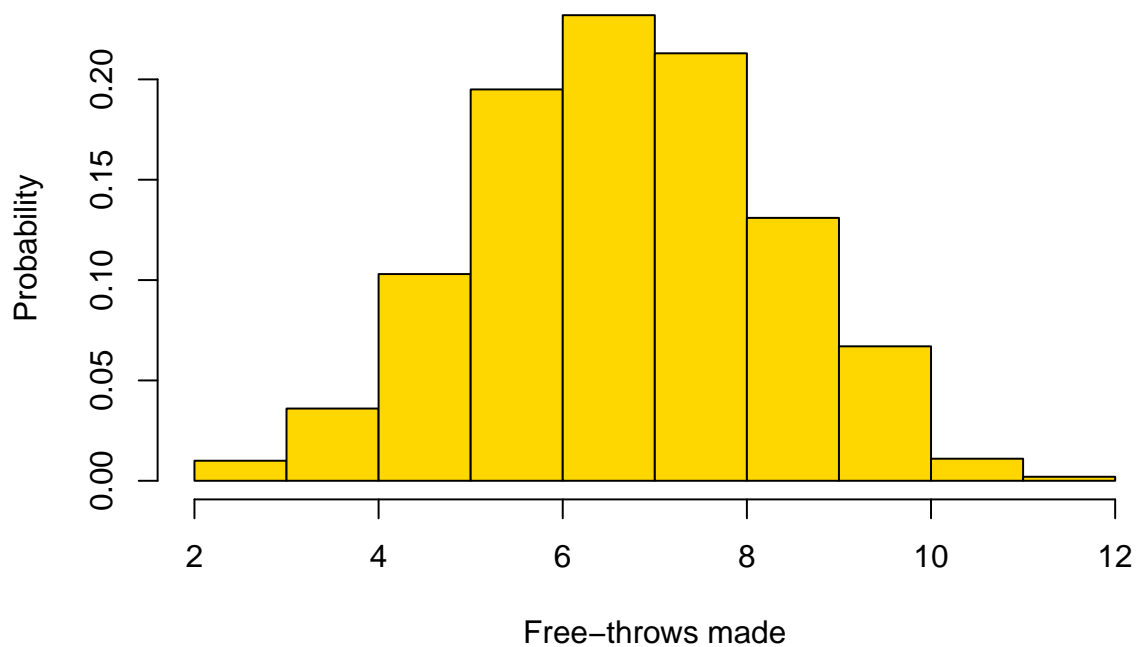
```
qbinom(p = 0.1, size = 12, prob = .6)
```

```
## [1] 5
```

Part D:

Plot the distributional histogram for the Bob's free-throw attempts by simulating this scenario.

```
x <- rbinom(n = 1000, size = 12, prob = .6)  
hist(x, probability = TRUE, main = "",  
     xlab = "Free-throws made",  
     ylab = "Probability", col = "Gold")
```

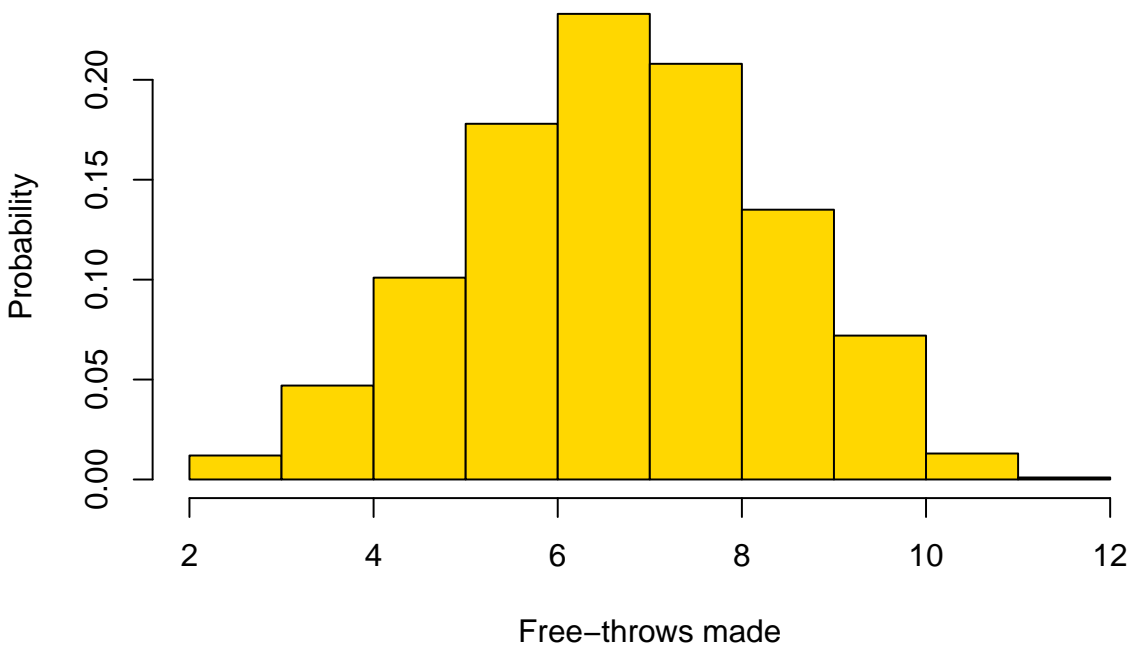


Part D:

Can you create this simulation design without using `rbinom()` function, instead `sample()` command.

```
x_sim <- numeric()
for(i in 1:1000){
  x_sim[i] <- sum(sample(0:1, size = 12, prob = c(0.4,0.6), replace = TRUE))
}

hist(x_sim, probability = TRUE, main = "",
      xlab = "Free-throws made",
      ylab = "Probability", col = "Gold")
```



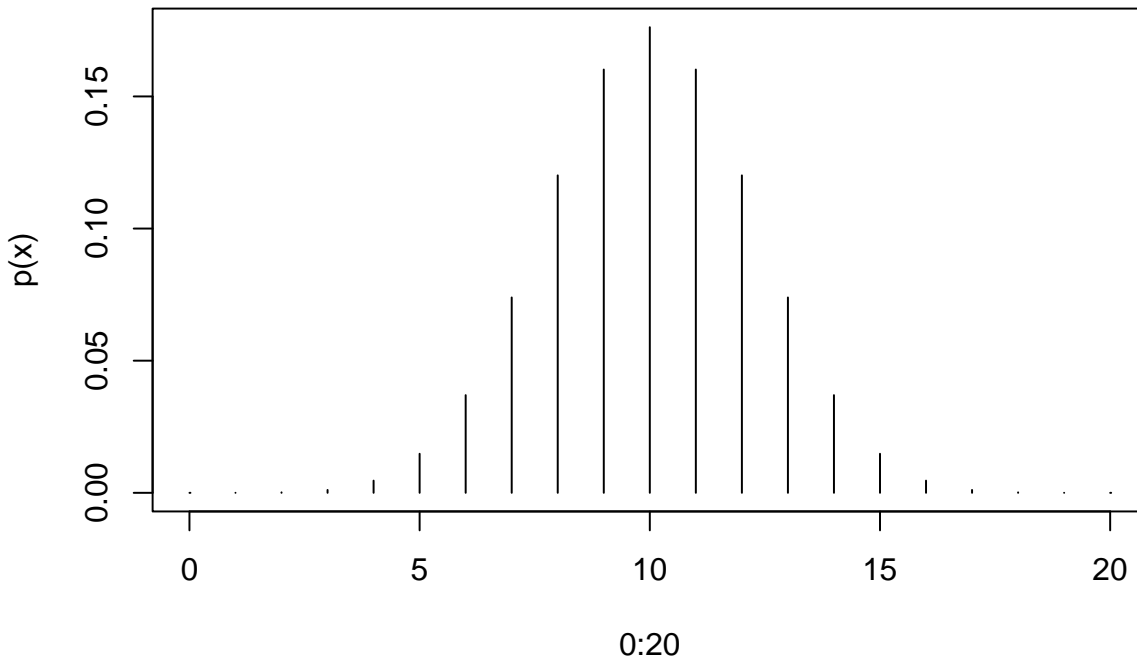
Exercise 2:

Dr. Horrible decided to give his Econ 1 students a pop quiz on Advanced Quantum Mechanics. Since he isn't completely unreasonable he made the quiz True-or-False. Since they don't know any Quantum Physics, Dr. Horrible's students guess randomly on each of the 20 questions.

Part A:

An individual student's score on this quiz can be modeled as the realization of random variable. What random variable? Plot its pmf.

```
plot(0:20, dbinom(0:20, size = 20, prob = 0.5), type = 'h', ylab = 'p(x)')
```



Part B:

Suppose that a passing grade on the quiz is a 60%. What is the probability that a given student passes?

```
pbinom(11, size=20, prob=.5, lower.tail=FALSE)
```

```
## [1] 0.2517223
```

```
# Alternatively,
```

```
# sum(dbinom(12:20, size = 20, prob = 0.5))
```

Part C:

Suppose that anything over 90% is an A. What is the probability that an individual student gets an A?

```
pbinom(17, size = 20, prob = .5, lower.tail = FALSE)
```

```
## [1] 0.0002012253
```

```
# Alternatively,  
# sum(dbinom(18:20, size = 20, prob = 0.5))
```

Part D:

If the class has 250 students, approximately how many will pass the quiz? Approximately how many will get an A?

```
pbinom(11, size=20, prob=.5, lower.tail=FALSE) * 250
```

```
## [1] 62.93058
```

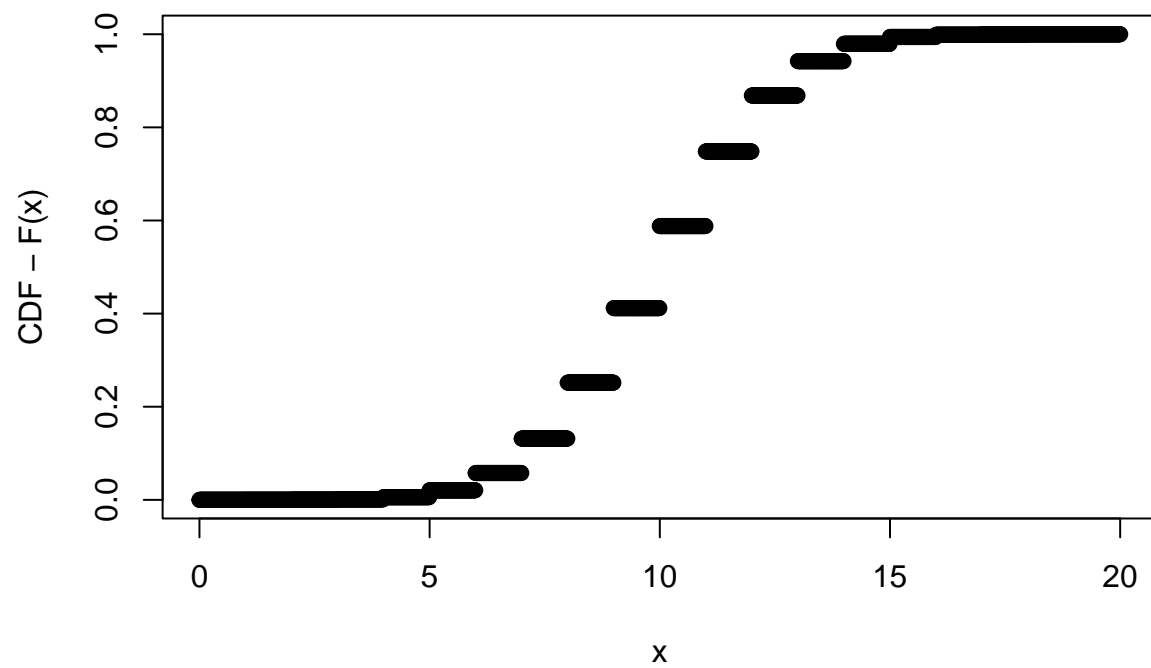
```
pbinom(17, size = 20, prob = .5, lower.tail = FALSE) * 250
```

```
## [1] 0.05030632
```

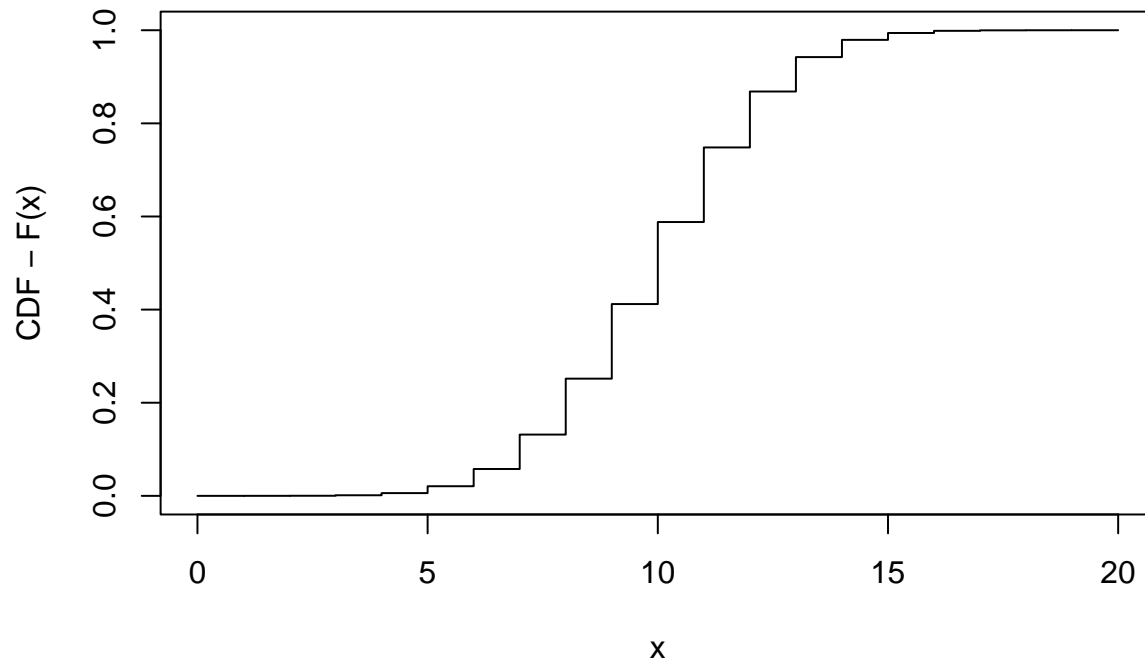
Part E:

Obtain a CDF plot for this distribution.

```
x <- seq(from = 0, to = 20, by = 0.01)  
y <- pbinom(x, size = 20, prob = .5)  
plot(x, y, ylab = 'CDF - F(x)')
```



```
plot(x, y, ylab = 'CDF - F(x)', type = "s")
```



Exercise 3:

Suppose that the data concerning the first-year salaries of METU graduates is normally distributed with the population mean $\mu = 60.000$ TL and the population standard deviation $\sigma = 15.000$ TL.

Part A:

Find the probability of a randomly selected METU graduate earning more than 80.000 TL annually.

```
pnorm(q = 80000, mean = 60000, sd = 15000, lower.tail = FALSE)
```

```
## [1] 0.09121122
```

Part B:

Find the lower bound of annual salaries of the top 15% earners of METU graduates.

```
qnorm(p = 0.15, mean = 60000, sd = 15000, lower.tail = FALSE)
```

```
## [1] 75546.5
```

Part C:

Find the probability of a randomly selected METU graduate earning between 50.000 and 55.000 TL annually.

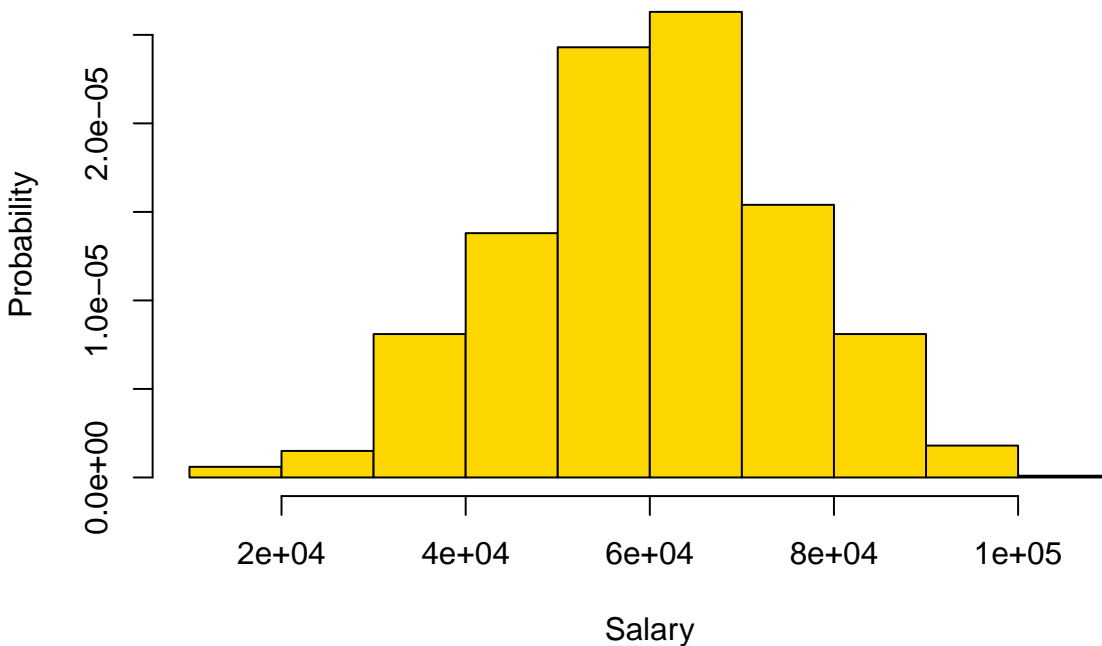
```
diff(pnorm(q = c(50000,55000), mean = 60000, sd = 15000))
```

```
## [1] 0.1169488
```

Part D:

Generate a new data set of size 1000, from normal distribution having the same parameters and then obtain a histogram.

```
x <- rnorm(n = 1000, mean = 60000, sd = 15000)
hist(x, probability = TRUE, main = "",
     xlab = "Salary",
     ylab = "Probability", col = "Gold")
```



Exceptions

Exercise 1:

Write an R function which calculates Euclidean distance of two given points. If the dimensions does not match then it stops.

```
distance <- function(u,v) {  
  if(length(u) == length(v)){  
  
    d <- sqrt(sum((u-v)^2))  
    return(round(d,3))  
  
  } else{  
  
    stop("vectors must have the same length")  
  }  
}
```

```
distance(u = 1:3, v = 3:9)
```

```
distance(u = c(2,2,-1), v = c(2,-1,5))
```

```
## [1] 6.708
```

Exercise 2:

Write a function that takes a real number (or a vector of numbers) x and a nonnegative integer n, and returns the value of the sum

$$1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \cdots + \frac{x^n}{n!} = \sum_{i=0}^n \frac{x^i}{i!}$$

If the parameter x is a vector, the function's output should be a vector.

If the parameter n is negative, the function should display warning message “parameter n should not be negative” and takes absolute value of n then uses abs(n) for the calculations.

```
powerseries <- function(x, n){  
  
  if(n < 0){  
    warning(paste("Given value n =",n,  
                  "is not a positive number. \n  
                  abs(n) will be used \n"))  
    n <- abs(n)  
  }  
  out_power_series <- numeric()  
  
  for(i in 1:length(x)){
```

```

    out_power_series[i] <- sum(x[i]^(0:n) / factorial(0:n))
  }

  return(out_power_series)
}
powerseries(x=1, n=7)
powerseries(x=2, n=-1)

```

Timings

Part A:

Generate 1000000 random numbers from standard normal distribution and sort them in increasing order. Calculate the run time.

```

start <- Sys.time()

x <- sort(rnorm(n = 1000000))

end <- Sys.time()

end - start

## Time difference of 0.09773898 secs

# alternatively,
system.time({
  sort(rnorm(n = 1000000))
})

##      user  system elapsed
##    0.08    0.03    0.11

```

Part B:

Now, create a for loop with 50 iterations where each iteration sleeps system 0.1 seconds (Sys.sleep(0.1)) and have a Progress Bar to watch the progress of the loop.

```

n_iter <- 50 # Number of iterations of the loop

# Initializes the progress bar
pb <- txtProgressBar(min = 0,
                     max = n_iter,
                     style = 3,

```

```

width = 50,
char = "=")

## |

for(i in 1:n_iter) {

  #-----
  # Code to be executed
  #-----

  Sys.sleep(0.1) # Remove this line and add your code

  #-----

  # Sets the progress bar to the current state
  setTxtProgressBar(pb, i)
}

## | |=

close(pb) # Close the connection

# alternatively

n_iter <- 50 # Number of iterations

pb <- winProgressBar(title = "Windows progress bar", # Window title
  label = "Percentage completed", # Window label
  min = 0, # Minimum value of the bar
  max = n_iter, # Maximum value of the bar
  initial = 0, # Initial value of the bar
  width = 300L) # Width of the window

for(i in 1:n_iter) {

  Sys.sleep(0.1) # Remove this line and add your code

  pctg <- paste(round(i/n_iter * 100, 0), "% completed")
  setWinProgressBar(pb, i, label = pctg)
}

beepr::beep(5)

```

```
close(pb) # Close the connection
```

```
## NULL
```